

# DIABETES PREDICTION

- Algorithm: Random Forest Classifier, K-Nearest Neighbors (KNN), adaboosting classifier, linear Regression, Decision Tree Classifier,
- Description: Predict whether a person has diabetes or not using features like glucose levels and BMI.

```
In [187]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import warnings
warnings.filterwarnings('ignore')
```

```
In [188]: df=pd.read_csv('diabetes_prediction_dataset.csv')
df
```

```
Out[188]:
```

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level	diabetes
0	Female	80.0	0	1	never	25.19	6.6	140	0
1	Female	54.0	0	0	No Info	27.32	6.6	80	0
2	Male	28.0	0	0	never	27.32	5.7	158	0
3	Female	36.0	0	0	current	23.45	5.0	155	0
4	Male	76.0	1	1	current	20.14	4.8	155	0
...	...	...	...	...	...	...	...	...	...
99995	Female	80.0	0	0	No Info	27.32	6.2	90	0
99996	Female	2.0	0	0	No Info	17.37	6.5	100	0
99997	Male	66.0	0	0	former	27.83	5.7	155	0
99998	Female	24.0	0	0	never	35.42	4.0	100	0
99999	Female	57.0	0	0	current	22.43	6.6	90	0

100000 rows × 9 columns

```
In [189]: df.head()
```

```
Out[189]:
```

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level	diabetes
0	Female	80.0	0	1	never	25.19	6.6	140	0
1	Female	54.0	0	0	No Info	27.32	6.6	80	0
2	Male	28.0	0	0	never	27.32	5.7	158	0
3	Female	36.0	0	0	current	23.45	5.0	155	0
4	Male	76.0	1	1	current	20.14	4.8	155	0

```
In [190]: df.tail()
```

```
Out[190]:
```

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level	diabetes
99995	Female	80.0	0	0	No Info	27.32	6.2	90	0
99996	Female	2.0	0	0	No Info	17.37	6.5	100	0
99997	Male	66.0	0	0	former	27.83	5.7	155	0
99998	Female	24.0	0	0	never	35.42	4.0	100	0
99999	Female	57.0	0	0	current	22.43	6.6	90	0

```
In [191]: df.isnull().sum()
```

```
Out[191]: gender                0
          age                  0
          hypertension          0
          heart_disease         0
          smoking_history       0
          bmi                   0
          HbA1c_level           0
          blood_glucose_level   0
          diabetes              0
          dtype: int64
```

```
In [192]: print("Number of rows =",df.shape[0])
          print("Number of columns =",df.shape[1])

Number of rows = 100000
Number of columns = 9
```

```
In [193]: df.size

Out[193]: 900000
```

```
In [194]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   gender                100000 non-null  object
1   age                   100000 non-null  float64
2   hypertension          100000 non-null  int64
3   heart_disease         100000 non-null  int64
4   smoking_history       100000 non-null  object
5   bmi                   100000 non-null  float64
6   HbA1c_level           100000 non-null  float64
7   blood_glucose_level   100000 non-null  int64
8   diabetes              100000 non-null  int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB
```

```
In [195]: df.describe()
```

```
Out[195]:
```


	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level	diabetes
count	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000
mean	41.885856	0.07485	0.039420	27.320767	5.527507	138.058060	0.085000
std	22.516840	0.26315	0.194593	6.636783	1.070672	40.708136	0.278883
min	0.080000	0.00000	0.000000	10.010000	3.500000	80.000000	0.000000
25%	24.000000	0.00000	0.000000	23.630000	4.800000	100.000000	0.000000
50%	43.000000	0.00000	0.000000	27.320000	5.800000	140.000000	0.000000
75%	60.000000	0.00000	0.000000	29.580000	6.200000	159.000000	0.000000
max	80.000000	1.00000	1.000000	95.690000	9.000000	300.000000	1.000000

```
In [196]: df['gender'].value_counts()
```

```
Out[196]: gender
          Female    58552
          Male     41430
          Other       18
          Name: count, dtype: int64
```

```
In [197]: df['gender'].value_counts().plot(kind='pie')

Out[197]: <Axes: ylabel='count'>
```

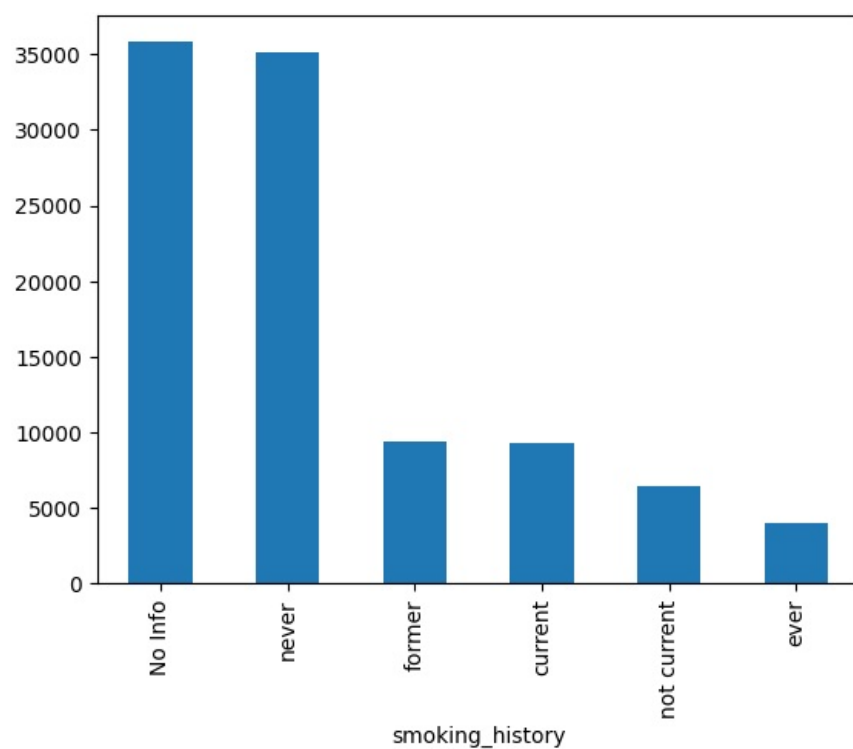


```
In [198]: df['smoking_history'].value_counts()
```

```
Out[198]: smoking_history
No Info      35816
never        35095
former        9352
current       9286
not current   6447
ever          4004
Name: count, dtype: int64
```

```
In [199]: df['smoking_history'].value_counts().plot(kind='bar')
```

```
Out[199]: <Axes: xlabel='smoking_history'>
```



In [200] df.columns

Out[200]: Index(['gender', 'age', 'hypertension', 'heart\_disease', 'smoking\_history', 'bmi', 'HbA1c\_level', 'blood\_glucose\_level', 'diabetes'], dtype='object')

In [201] *#Convert categorical feature into numerical feature by using label encoding technique and other technique using*  
 encode = LabelEncoder()  
 df['smoking\_history']=encode.fit\_transform(df['smoking\_history'])  
 df['gender']=np.where(df['gender'].str.contains('Female'),0,1)

In [202] df.head()

Out[202]:

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level	diabetes
0	0	80.0	0	1	4	25.19	6.6	140	0
1	0	54.0	0	0	0	27.32	6.6	80	0
2	1	28.0	0	0	4	27.32	5.7	158	0
3	0	36.0	0	0	1	23.45	5.0	155	0
4	1	76.0	1	1	1	20.14	4.8	155	0

In [203] *#Take relevant features from the dataset to train our model.*  
 x = df[['smoking\_history', 'gender', 'bmi', 'blood\_glucose\_level']]

```
y = df['diabetes']
```

```
In [204... #Doing standard scaling
scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
```

```
In [205... x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size=0.2, random_state=42)
```

```
In [206... # Define models
models = [
    ("Random Forest", RandomForestClassifier(random_state=42)),
    ("K-Nearest Neighbors", KNeighborsClassifier()),
    ("AdaBoost Classifier", AdaBoostClassifier(random_state=42)),
    ("Decision Tree Classifier", DecisionTreeClassifier(random_state=42)),
    ("Logistic Regression", LogisticRegression(random_state=42))
]
```

```
In [207... # Train and evaluate models, storing results
results = {}
for name, model in models:
    model.fit(x_train, y_train)
    predictions = model.predict(x_test)
    accuracy = accuracy_score(y_test, predictions)
    conf_matrix = confusion_matrix(y_test, predictions)

    results[name] = {
        "Accuracy": accuracy,
        "Confusion Matrix": conf_matrix
    }

    print(f"{name} Accuracy: {accuracy:.4f}")
    print(classification_report(y_test, predictions, target_names=["Negative", "Positive"]))
    print("-" * 50)
```

Random Forest Accuracy: 0.9290					
	precision	recall	f1-score	support	
Negative	0.95	0.97	0.96	18292	
Positive	0.62	0.44	0.51	1708	
accuracy			0.93	20000	
macro avg	0.78	0.71	0.74	20000	
weighted avg	0.92	0.93	0.92	20000	

K-Nearest Neighbors Accuracy: 0.9431					
	precision	recall	f1-score	support	
Negative	0.95	0.99	0.97	18292	
Positive	0.85	0.41	0.55	1708	
accuracy			0.94	20000	
macro avg	0.90	0.70	0.76	20000	
weighted avg	0.94	0.94	0.93	20000	

AdaBoost Classifier Accuracy: 0.9473					
	precision	recall	f1-score	support	
Negative	0.95	1.00	0.97	18292	
Positive	1.00	0.38	0.55	1708	
accuracy			0.95	20000	
macro avg	0.97	0.69	0.76	20000	
weighted avg	0.95	0.95	0.94	20000	

Decision Tree Classifier Accuracy: 0.9241					
	precision	recall	f1-score	support	
Negative	0.95	0.97	0.96	18292	
Positive	0.57	0.47	0.51	1708	
accuracy			0.92	20000	
macro avg	0.76	0.72	0.74	20000	
weighted avg	0.92	0.92	0.92	20000	


Logistic Regression Accuracy: 0.9401					
	precision	recall	f1-score	support	
Negative	0.94	1.00	0.97	18292	
Positive	0.87	0.35	0.50	1708	
accuracy			0.94	20000	
macro avg	0.91	0.67	0.73	20000	
weighted avg	0.94	0.94	0.93	20000	

```
In [208]: # Plot the confusion matrices
fig, axes = plt.subplots(2, 2, figsize=(14, 14))
fig.suptitle('Confusion Matrices of Different Models', fontsize=20)

for (name, result), ax in zip(results.items(), axes.flatten()):
    sns.heatmap(result['Confusion Matrix'], annot=True, fmt='d', cmap='Blues', ax=ax)
    ax.set_title(f'{name} (Accuracy: {result["Accuracy"]:.4f})')
    ax.set_xlabel('Predicted')
    ax.set_ylabel('Actual')

plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()
```

## Confusion Matrices of Different Models



In [ ]:

## Take New dataset for testing the models

In [209..

```
new_sample_data = pd.DataFrame({
    'smoking_history': np.random.choice([0, 1, 2], 10), # 0: never, 1: current, 2: former
    'gender': np.random.choice([0, 1], 10), # 0: Male, 1: Female
    'bmi': np.random.uniform(18, 40, 10),
    'blood_glucose_level': np.random.uniform(70, 200, 10)
})
```

In [210..

new\_sample\_data

Out[210]:

	smoking_history	gender	bmi	blood_glucose_level
0	0	1	26.604753	70.075834
1	0	1	31.469352	141.021477
2	2	1	18.445357	82.479776
3	0	0	18.130357	156.660308
4	2	1	24.560052	151.634666
5	1	0	38.117731	140.558481
6	1	0	30.616946	181.308104
7	1	1	29.731554	117.696247
8	2	0	19.389524	166.858844
9	0	0	22.620052	71.649262

```
In [211... # Scale the new sample data
new_synthetic_data_scaled = scaler.transform(new_sample_data)
```

```
In [212... # Predictions on new sample data
new_predictions = {}
for name, model in models:
    new_predictions[name] = model.predict(new_synthetic_data_scaled)
    print(f"{name} Predictions on New Data:", new_predictions[name])
```

```
Random Forest Predictions on New Data: [0 0 0 0 1 0 0 0 0 0]
K-Nearest Neighbors Predictions on New Data: [0 0 0 0 0 0 0 0 0 0]
AdaBoost Classifier Predictions on New Data: [0 0 0 0 0 0 0 0 0 0]
Decision Tree Classifier Predictions on New Data: [0 0 0 0 0 0 1 0 0 0]
Logistic Regression Predictions on New Data: [0 0 0 0 0 0 0 0 0 0]
```

```
In [ ]:
```

## Important Note:-

Our problem statement specifies one important algorithm to use for solving this particular problem, which is linear regression. However, we are not using this algorithm. In our problem, where the target is to predict whether a person has diabetes (a binary classification problem), logistic regression is more appropriate than linear regression. That is why we use this linear regression algorithm.

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js