

HOUSE PRICE PREDICTION

- Algorithm: Linear Regression, random forest regression, decision tree regression, gradient boosting regressor
- Description: Predict house prices based on features like area, number of bedrooms, and location.
- For dataset-[here](#)

```
In [359]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import mean_squared_error
import warnings
warnings.filterwarnings('ignore')
```

```
In [360]: #loading the datasets
df=pd.read_csv('data.csv')
df
```

```
Out[360]:
```

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	sqft_above	sqft_basement	yr_built
0	2014-05-02 00:00:00	3.130000e+05	3.0	1.50	1340	7912	1.5	0	0	3	1340	0	
1	2014-05-02 00:00:00	2.384000e+06	5.0	2.50	3650	9050	2.0	0	4	5	3370	280	
2	2014-05-02 00:00:00	3.420000e+05	3.0	2.00	1930	11947	1.0	0	0	4	1930	0	
3	2014-05-02 00:00:00	4.200000e+05	3.0	2.25	2000	8030	1.0	0	0	4	1000	1000	
4	2014-05-02 00:00:00	5.500000e+05	4.0	2.50	1940	10500	1.0	0	0	4	1140	800	
...
4595	2014-07-09 00:00:00	3.081667e+05	3.0	1.75	1510	6360	1.0	0	0	4	1510	0	
4596	2014-07-09 00:00:00	5.343333e+05	3.0	2.50	1460	7573	2.0	0	0	3	1460	0	
4597	2014-07-09 00:00:00	4.169042e+05	3.0	2.50	3010	7014	2.0	0	0	3	3010	0	
4598	2014-07-10 00:00:00	2.034000e+05	4.0	2.00	2090	6630	1.0	0	0	3	1070	1020	
4599	2014-07-10 00:00:00	2.206000e+05	3.0	2.50	1490	8102	2.0	0	0	4	1490	0	

4600 rows × 18 columns

```
In [361]: df.head()
```

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	sqft_above	sqft_basement	yr_built
0	2014-05-02 00:00:00	313000.0	3.0	1.50	1340	7912	1.5	0	0	3	1340	0	1955
1	2014-05-02 00:00:00	2384000.0	5.0	2.50	3650	9050	2.0	0	4	5	3370	280	1921
2	2014-05-02 00:00:00	342000.0	3.0	2.00	1930	11947	1.0	0	0	4	1930	0	1966
3	2014-05-02 00:00:00	420000.0	3.0	2.25	2000	8030	1.0	0	0	4	1000	1000	1963
4	2014-05-02 00:00:00	550000.0	4.0	2.50	1940	10500	1.0	0	0	4	1140	800	1976

In [362]: `df.tail()`

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	sqft_above	sqft_basement	y
4595	2014-07-09 00:00:00	308166.666667	3.0	1.75	1510	6360	1.0	0	0	4	1510	0	
4596	2014-07-09 00:00:00	534333.333333	3.0	2.50	1460	7573	2.0	0	0	3	1460	0	
4597	2014-07-09 00:00:00	416904.166667	3.0	2.50	3010	7014	2.0	0	0	3	3010	0	
4598	2014-07-10 00:00:00	203400.000000	4.0	2.00	2090	6630	1.0	0	0	3	1070	1020	
4599	2014-07-10 00:00:00	220600.000000	3.0	2.50	1490	8102	2.0	0	0	4	1490	0	

In [363]: `print("The Number of rows",df.shape[0])`
`print("The Number of columns",df.shape[1])`

The Number of rows 4600
The Number of columns 18

In [364]: `df.size`

Out[364]: 82800

In [365]: `#lets checks the numerical and categorical feature`
`Categorical_feature=[feature for feature in df.columns if df[feature].dtype=='0']`
`Numerical_feature=[feature for feature in df.columns if df[feature].dtype!='0']`
`print("The Categorical feature is =",Categorical_feature)`
`print("The Numerical feature is =",Numerical_feature)`
The Categorical feature is = ['date', 'street', 'city', 'statezip', 'country']
The Numerical feature is = ['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated']

In [366]: `df.isnull().sum()`

Out[366]:

date	0
price	0
bedrooms	0
bathrooms	0
sqft_living	0
sqft_lot	0
floors	0
waterfront	0
view	0
condition	0
sqft_above	0
sqft_basement	0
yr_built	0
yr_renovated	0
street	0
city	0
statezip	0
country	0
dtype: int64	

In [367]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4600 entries, 0 to 4599
Data columns (total 18 columns):
#   Column              Non-Null Count  Dtype
---  -
0   date                4600 non-null   object
1   price               4600 non-null   float64
2   bedrooms            4600 non-null   float64
3   bathrooms            4600 non-null   float64
4   sqft_living          4600 non-null   int64
5   sqft_lot             4600 non-null   int64
6   floors              4600 non-null   float64
7   waterfront           4600 non-null   int64
8   view                4600 non-null   int64
9   condition            4600 non-null   int64
10  sqft_above           4600 non-null   int64
11  sqft_basement        4600 non-null   int64
12  yr_built             4600 non-null   int64
13  yr_renovated         4600 non-null   int64
14  street               4600 non-null   object
15  city                 4600 non-null   object
16  statezip             4600 non-null   object
17  country              4600 non-null   object
dtypes: float64(4), int64(9), object(5)
memory usage: 647.0+ KB

```

```
In [368]: print(df.duplicated().sum())
```

```
0
```

```
In [369]: df.describe()
```

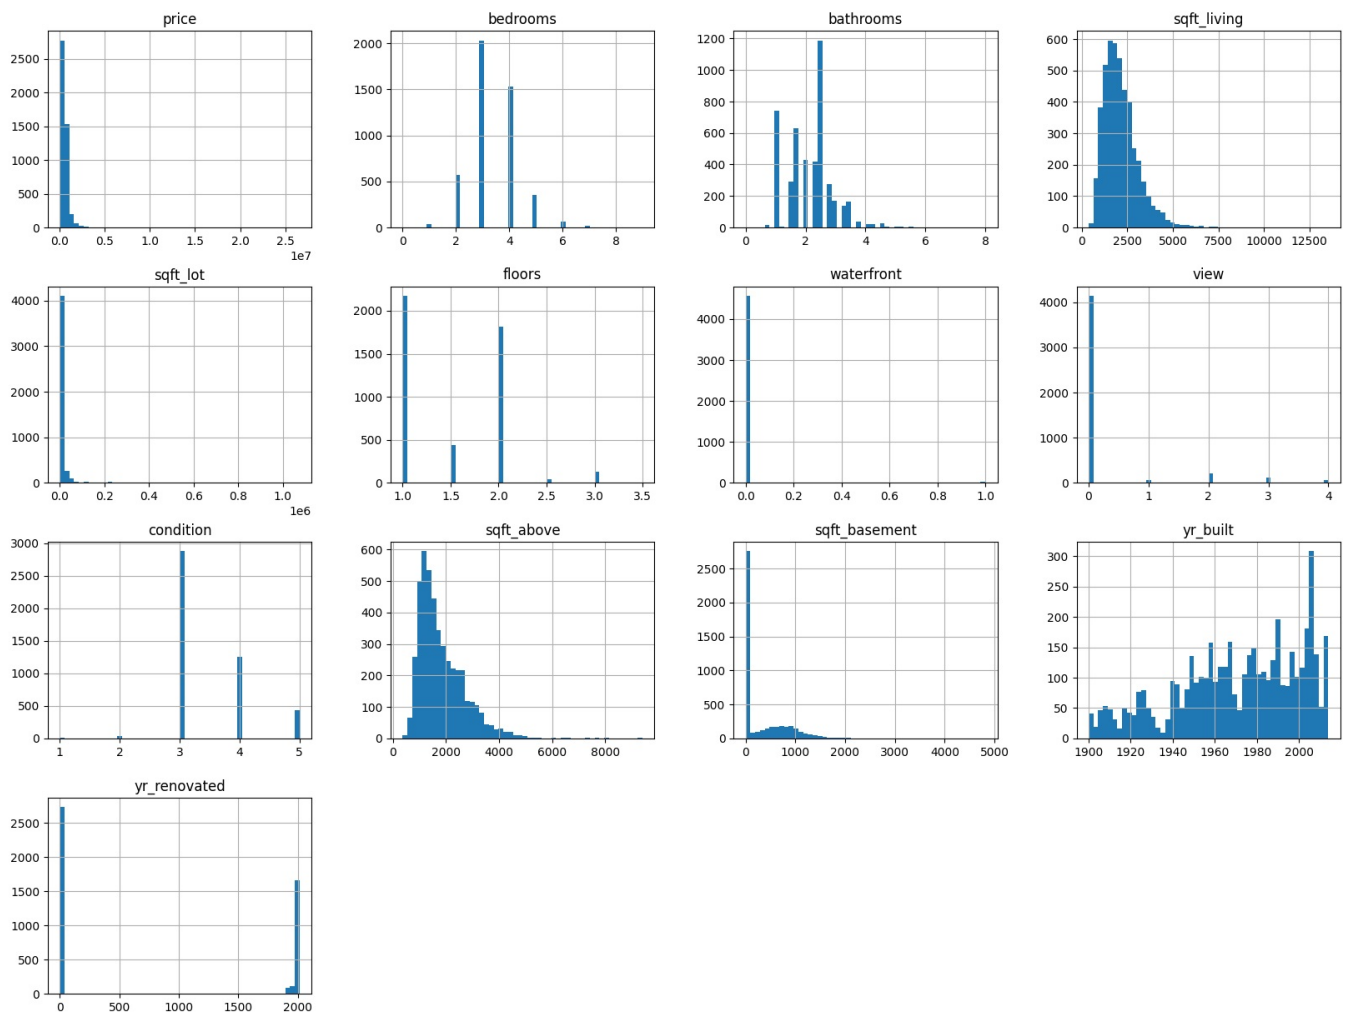
```
Out[369]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	sqft_abov
count	4.600000e+03	4600.000000	4600.000000	4600.000000	4.600000e+03	4600.000000	4600.000000	4600.000000	4600.000000	4600.000000
mean	5.519630e+05	3.400870	2.160815	2139.346957	1.485252e+04	1.512065	0.007174	0.240652	3.451739	1827.2654
std	5.638347e+05	0.908848	0.783781	963.206916	3.588444e+04	0.538288	0.084404	0.778405	0.677230	862.1689
min	0.000000e+00	0.000000	0.000000	370.000000	6.380000e+02	1.000000	0.000000	0.000000	1.000000	370.0000
25%	3.228750e+05	3.000000	1.750000	1460.000000	5.000750e+03	1.000000	0.000000	0.000000	3.000000	1190.0000
50%	4.609435e+05	3.000000	2.250000	1980.000000	7.683000e+03	1.500000	0.000000	0.000000	3.000000	1590.0000
75%	6.549625e+05	4.000000	2.500000	2620.000000	1.100125e+04	2.000000	0.000000	0.000000	4.000000	2300.0000
max	2.659000e+07	9.000000	8.000000	13540.000000	1.074218e+06	3.500000	1.000000	4.000000	5.000000	9410.0000


```
In [370]: df.columns
```

```
Out[370]: Index(['date', 'price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot',
        'floors', 'waterfront', 'view', 'condition', 'sqft_above',
        'sqft_basement', 'yr_built', 'yr_renovated', 'street', 'city',
        'statezip', 'country'],
        dtype='object')
```

```
In [371]: df.hist(bins=50, figsize=(20,15))
plt.show()
```

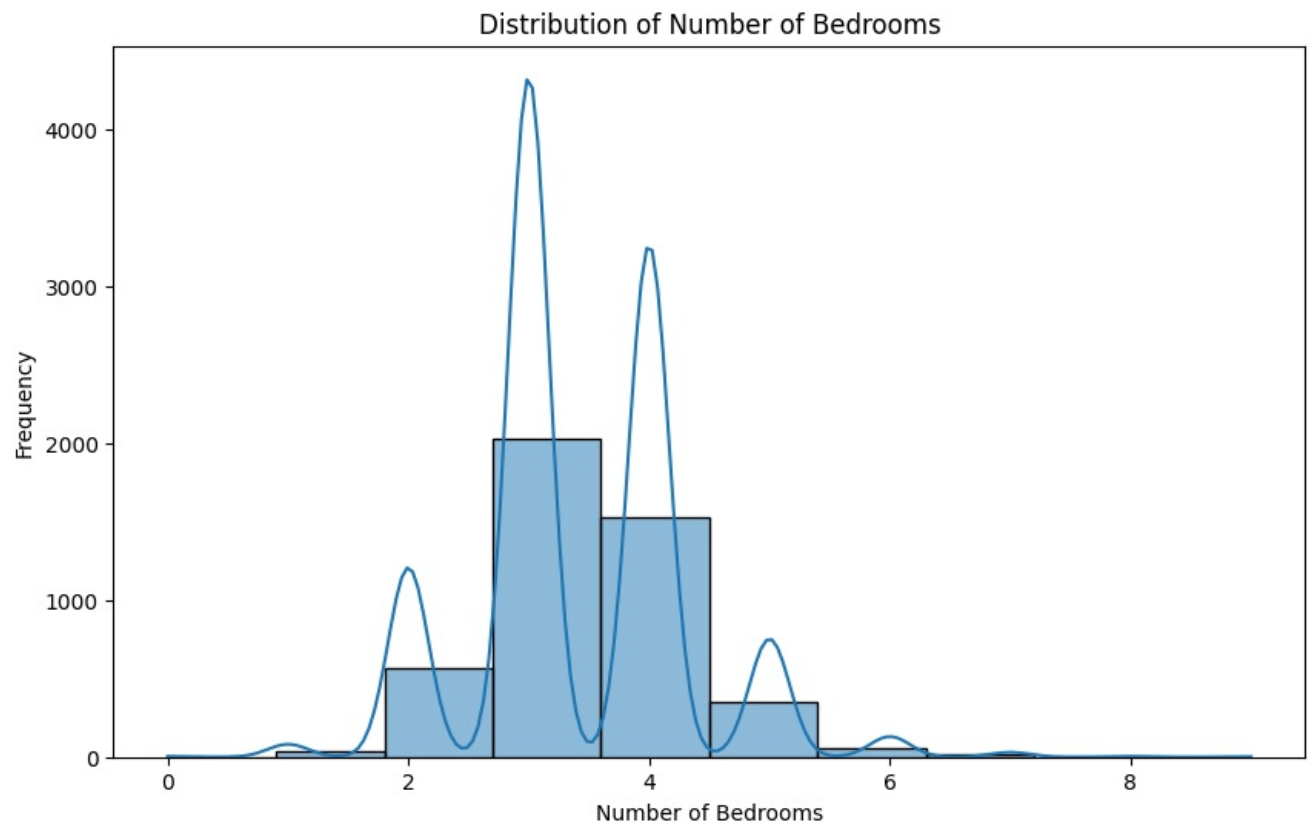


```
In [372... # Visualize the distribution of house prices
plt.figure(figsize=(10, 6))
sns.histplot(df['price'], kde=True)
plt.title('Distribution of House Prices')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()
```




```
In [373... # Visualize the distribution of number of bedrooms
plt.figure(figsize=(10, 6))
sns.histplot(df['bedrooms'], kde=True, bins=df['bedrooms'].nunique())
```

```
plt.title('Distribution of Number of Bedrooms')
plt.xlabel('Number of Bedrooms')
plt.ylabel('Frequency')
plt.show()
```



```
In [374]: # Visualize the distribution of cities using barplot
city_counts = df['city'].value_counts()
plt.figure(figsize=(20,13))
sns.barplot(x=city_counts.index, y=city_counts.values)
plt.title('Distribution of House Cities')
plt.xlabel('City')
plt.ylabel('Frequency')
plt.xticks(rotation=45)
plt.show()
```



```
In [375]: df.keys()
```

```
Out[375]: Index(['date', 'price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot',  
        'floors', 'waterfront', 'view', 'condition', 'sqft_above',  
        'sqft_basement', 'yr_built', 'yr_renovated', 'street', 'city',  
        'statezip', 'country'],  
        dtype='object')
```

```
In [376]: label = LabelEncoder()  
df['city']=label.fit_transform(df['city'])
```

```
In [377]: # Feature selection  
x = df[["bedrooms", "bathrooms", "city"]]  
y = df['price']
```

```
In [378]: x
```

```
Out[378]:
```

	bedrooms	bathrooms	city
0	3.0	1.50	36
1	5.0	2.50	35
2	3.0	2.00	18
3	3.0	2.25	3
4	4.0	2.50	31
...
4595	3.0	1.75	35
4596	3.0	2.50	3
4597	3.0	2.50	32
4598	4.0	2.00	35
4599	3.0	2.50	9

4600 rows × 3 columns

```
In [379]: y
```

```
Out[379]:
```

0	3.1300000e+05
1	2.3840000e+06
2	3.4200000e+05
3	4.2000000e+05
4	5.5000000e+05
...	...
4595	3.081667e+05
4596	5.343333e+05
4597	4.169042e+05
4598	2.034000e+05
4599	2.206000e+05

Name: price, Length: 4600, dtype: float64

```
In [380]: # Split the dataset into training and testing sets  
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
In [381]: #scaling  
scale = StandardScaler()  
x_train = scale.fit_transform(x_train)  
x_test = scale.transform(x_test)
```


```
In [382]: # Initialize models  
models = {  
    'Linear Regression': LinearRegression(),  
    'Random Forest': RandomForestRegressor(random_state=42),  
    'Decision Tree': DecisionTreeRegressor(random_state=42),  
    'Gradient Boosting': GradientBoostingRegressor(random_state=42)  
}
```

```
In [383]: # Train and evaluate models  
results = {}  
for name, model in models.items():  
    model.fit(x_train, y_train)  
    y_pred = model.predict(x_test)  
    mse = mean_squared_error(y_test, y_pred)  
    results[name] = mse  
    print(f'{name} Mean Squared Error: {mse}')
```

Linear Regression Mean Squared Error: 994739551125.2224
Random Forest Mean Squared Error: 990153673453.1986
Decision Tree Mean Squared Error: 1001240109347.2585
Gradient Boosting Mean Squared Error: 989400668255.3132

```
In [384]: # Create a DataFrame for the results  
results_df = pd.DataFrame(list(results.items()), columns=['Model', 'Mean Squared Error'])
```


```
In [385... # Visualize the results using a heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(results_df.set_index('Model').T, annot=True, cmap='coolwarm')
plt.title('Mean Squared Error of Different Models')
plt.xlabel('Model')
plt.ylabel('Mean Squared Error')
plt.show()
```



```
In [386... # Visualize the results
model_names = list(results.keys())
mse_values = list(results.values())
print(model_names)
print(mse_values)

['Linear Regression', 'Random Forest', 'Decision Tree', 'Gradient Boosting']
[np.float64(994739551125.2224), np.float64(990153673453.1986), np.float64(1001240109347.2585), np.float64(989400668255.3132)]
```

```
In [387... plt.figure(figsize=(10, 6))
sns.barplot(x=model_names, y=mse_values)
plt.title('Mean Squared Error of Different Models')
plt.xlabel('Model')
plt.ylabel('Mean Squared Error')
plt.show()
```



```
In [388]: #Take new datasets for prediction purpose
df=pd.read_csv('output.csv')
df.head()
```

```
Out[388]:
```

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	sqft_above	sqft_basement	yr_built
0	2014-05-02 00:00:00	313000.0	3.0	1.50	1340	7912	1.5	0	0	3	1340	0	1955
1	2014-05-02 00:00:00	2384000.0	5.0	2.50	3650	9050	2.0	0	4	5	3370	280	1921
2	2014-05-02 00:00:00	342000.0	3.0	2.00	1930	11947	1.0	0	0	4	1930	0	1966
3	2014-05-02 00:00:00	420000.0	3.0	2.25	2000	8030	1.0	0	0	4	1000	1000	1963
4	2014-05-02 00:00:00	550000.0	4.0	2.50	1940	10500	1.0	0	0	4	1140	800	1976

```
In [389]: label1 = LabelEncoder()
df['city']=label1.fit_transform(df['city'])
```

```
In [390]: df.keys()
```

```
Out[390]: Index(['date', 'price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot',
        'floors', 'waterfront', 'view', 'condition', 'sqft_above',
        'sqft_basement', 'yr_built', 'yr_renovated', 'street', 'city',
        'statezip', 'country'],
        dtype='object')
```

```
In [391]: new_df=df[["bedrooms","bathrooms","city"]]
new_df
```


Out[391]:

	bedrooms	bathrooms	city
0	3.0	1.50	36
1	5.0	2.50	35
2	3.0	2.00	18
3	3.0	2.25	3
4	4.0	2.50	31
...
4595	3.0	1.75	35
4596	3.0	2.50	3
4597	3.0	2.50	32
4598	4.0	2.00	35
4599	3.0	2.50	9

4600 rows × 3 columns

In [392...

```
# Scale the new sample data
test_data = scale.transform(new_df)
```

In [393...

```
# Predictions on new sample data
new_predictions = {}
for name, model in models.items():
    new_predictions[name] = model.predict(test_data)
print(f"{name} Predictions on New Data:", new_predictions[name])
```

Linear Regression Predictions on New Data: [410373.01969841 691687.87750807 477770.37191101 ... 628418.92546631 549729.64858226 568582.12450125]
Random Forest Predictions on New Data: [331453.3424912 726480.39768855 296844.940275 ... 399556.33280588 551040.55813699 278664.60714286]
Decision Tree Predictions on New Data: [327000. 714807.69230769 293664.28571429 ... 400055.96666668 547474.57763974 275750.]
Gradient Boosting Predictions on New Data: [359418.35056968 680980.46587945 302807.88126357 ... 482004.18226315 575506.47496632 408843.43443016]

In []:

In []:

In []: